



Home of Redis

Durability

Data persistence

Takeaways

Redis is an in-memory database. In a nutshell:

1. "NoSQL" is a way to think about managing data
2. All data is always written to and read from RAM^e
3. Data may be persisted to and recovered from durable storage

^e Redis^e Flash uses Flash as a RAM extension

What's fast, expensive and volatile?

(hint: not the SpaceX rocket)

What's fast, expensive and volatile?

RAM, which Redis uses to store all data:

- The fastest storage available
- Supports random access
- Provides similar latencies for both reads and writes

But RAM:

- Is expensive compared to SSD and spinning metal
- Loses data without power

Redis as a cache

Because RAM is volatile, using Redis as a cache for transient data is popular use case.

Types of volatile information include:

- Data copied from a different source of truth, e.g. DB
- Precomputed values, e.g. rendered HTML
- Ephemeral data, e.g. session store

Redis-as-a-Cache

- Data persistence is optional and tunable
- Configurable data memory footprint (`maxmemory`)
- Volatile keys that expire after a settable TTL
- OOM eviction policies (`maxmemory-policy`)
 - `allkeys-lru`, `allkeys-lfu`, `allkeys-random`, `noeviction`
 - `volatile-lru`, `volatile-lfu`, `volatile-random`, `volatile-ttl`
 - Are approximated for performance (`maxmemory-samples`)

transient data != availability

While the data in the cache isn't "important", ensuring its availability usually is.

- No cache means increase in application latency
- Cold cache means it needs to be warmed
- Beware of thundering herds toppling your infrastructure
- Life happens - recovery could take a while...

Redis as a primary database

(i.e. the data in it exists nowhere else)

Redis as a primary database

For some data Redis is the single source of truth.

Popular common examples are job queues, session metrics, stream analytics...

But even entire apps, for example [Muut](#) or [Spot.IM](#)

You really don't want to loose the data then.


Data durability

- Durability is the ability to recover the data after a failure, e.g. a power outage
- RAM is volatile, which makes it the undurable
- SSD and spinning metal disks do persist data
- Redis has two independent data persistence options: RDB and AOF
- Redis data persistence is 🐟-able:
 - Think different
 - Persistence impacts performance - costs CPU, RAM and IO

Dump files, a.k.a. RDB

- The dump file contains a serialization of all data
- Essentially a snapshot in time
- Can be created with SAVE (blocking!) or BGSAVE (not)
- Configurable background jobs with the save directive:
 - Syntax: `save <seconds> <changes>`
 - Defaults: `save 900 1`
`save 300 10`
`save 60 10000`
- Snapshots are atomic

BG: thread forking, RAM & Copy-on-Write



Append-only files

- The AOF is a log, write commands are added to it
- Enabled by setting `appendonly yes`
- Flush to disk configurable via `appendfsync`:
 - Every command (slow, not supported by Redis^e): `always`
 - Every second (default): `everysec`
 - The OS' consideration: `none`
- Can grow indefinitely, but manageable with:
 - Manual call to [BGREWRITEAOF](#)
 - `auto-aof-rewrite-percentage 100`
 - `auto-aof-rewrite-min-size 64mb`

RDB and/or AOF

- RDB files are
 - Compact in size and require minimal time to load
 - Atomic
 - Good for backups and disaster recovery
- AOF files are
 - Potentially large and time consuming to load
 - Asynchronously written
 - For making recent changes durable
- Use one, the other or both according to requirements

persisted data != availability

Ensuring that data is recoverable from disk after failure increases, but does not guarantee, availability.

- Reading serialized data from disk takes recovery time
- Replaying logged commands takes recovery time
- During that period the database isn't available
- That's what high-availability (replication and failover) is for