



Redis and Redis^e Architecture

Cihan B. - VP of Product Management

Agenda

Redis and Redis^e (Redis Enterprise) Architecture

- Shards, Nodes, Clusters
- Replication in Redis^e
- Anatomy of Read/Write Operation

Scaling Applications with Redis^e

- Scaling Throughput – Redis^e Resharding and Rebalancing
- Scaling Connections - Redis^e Proxy
- Scaling Data Size – Redis^e Flash

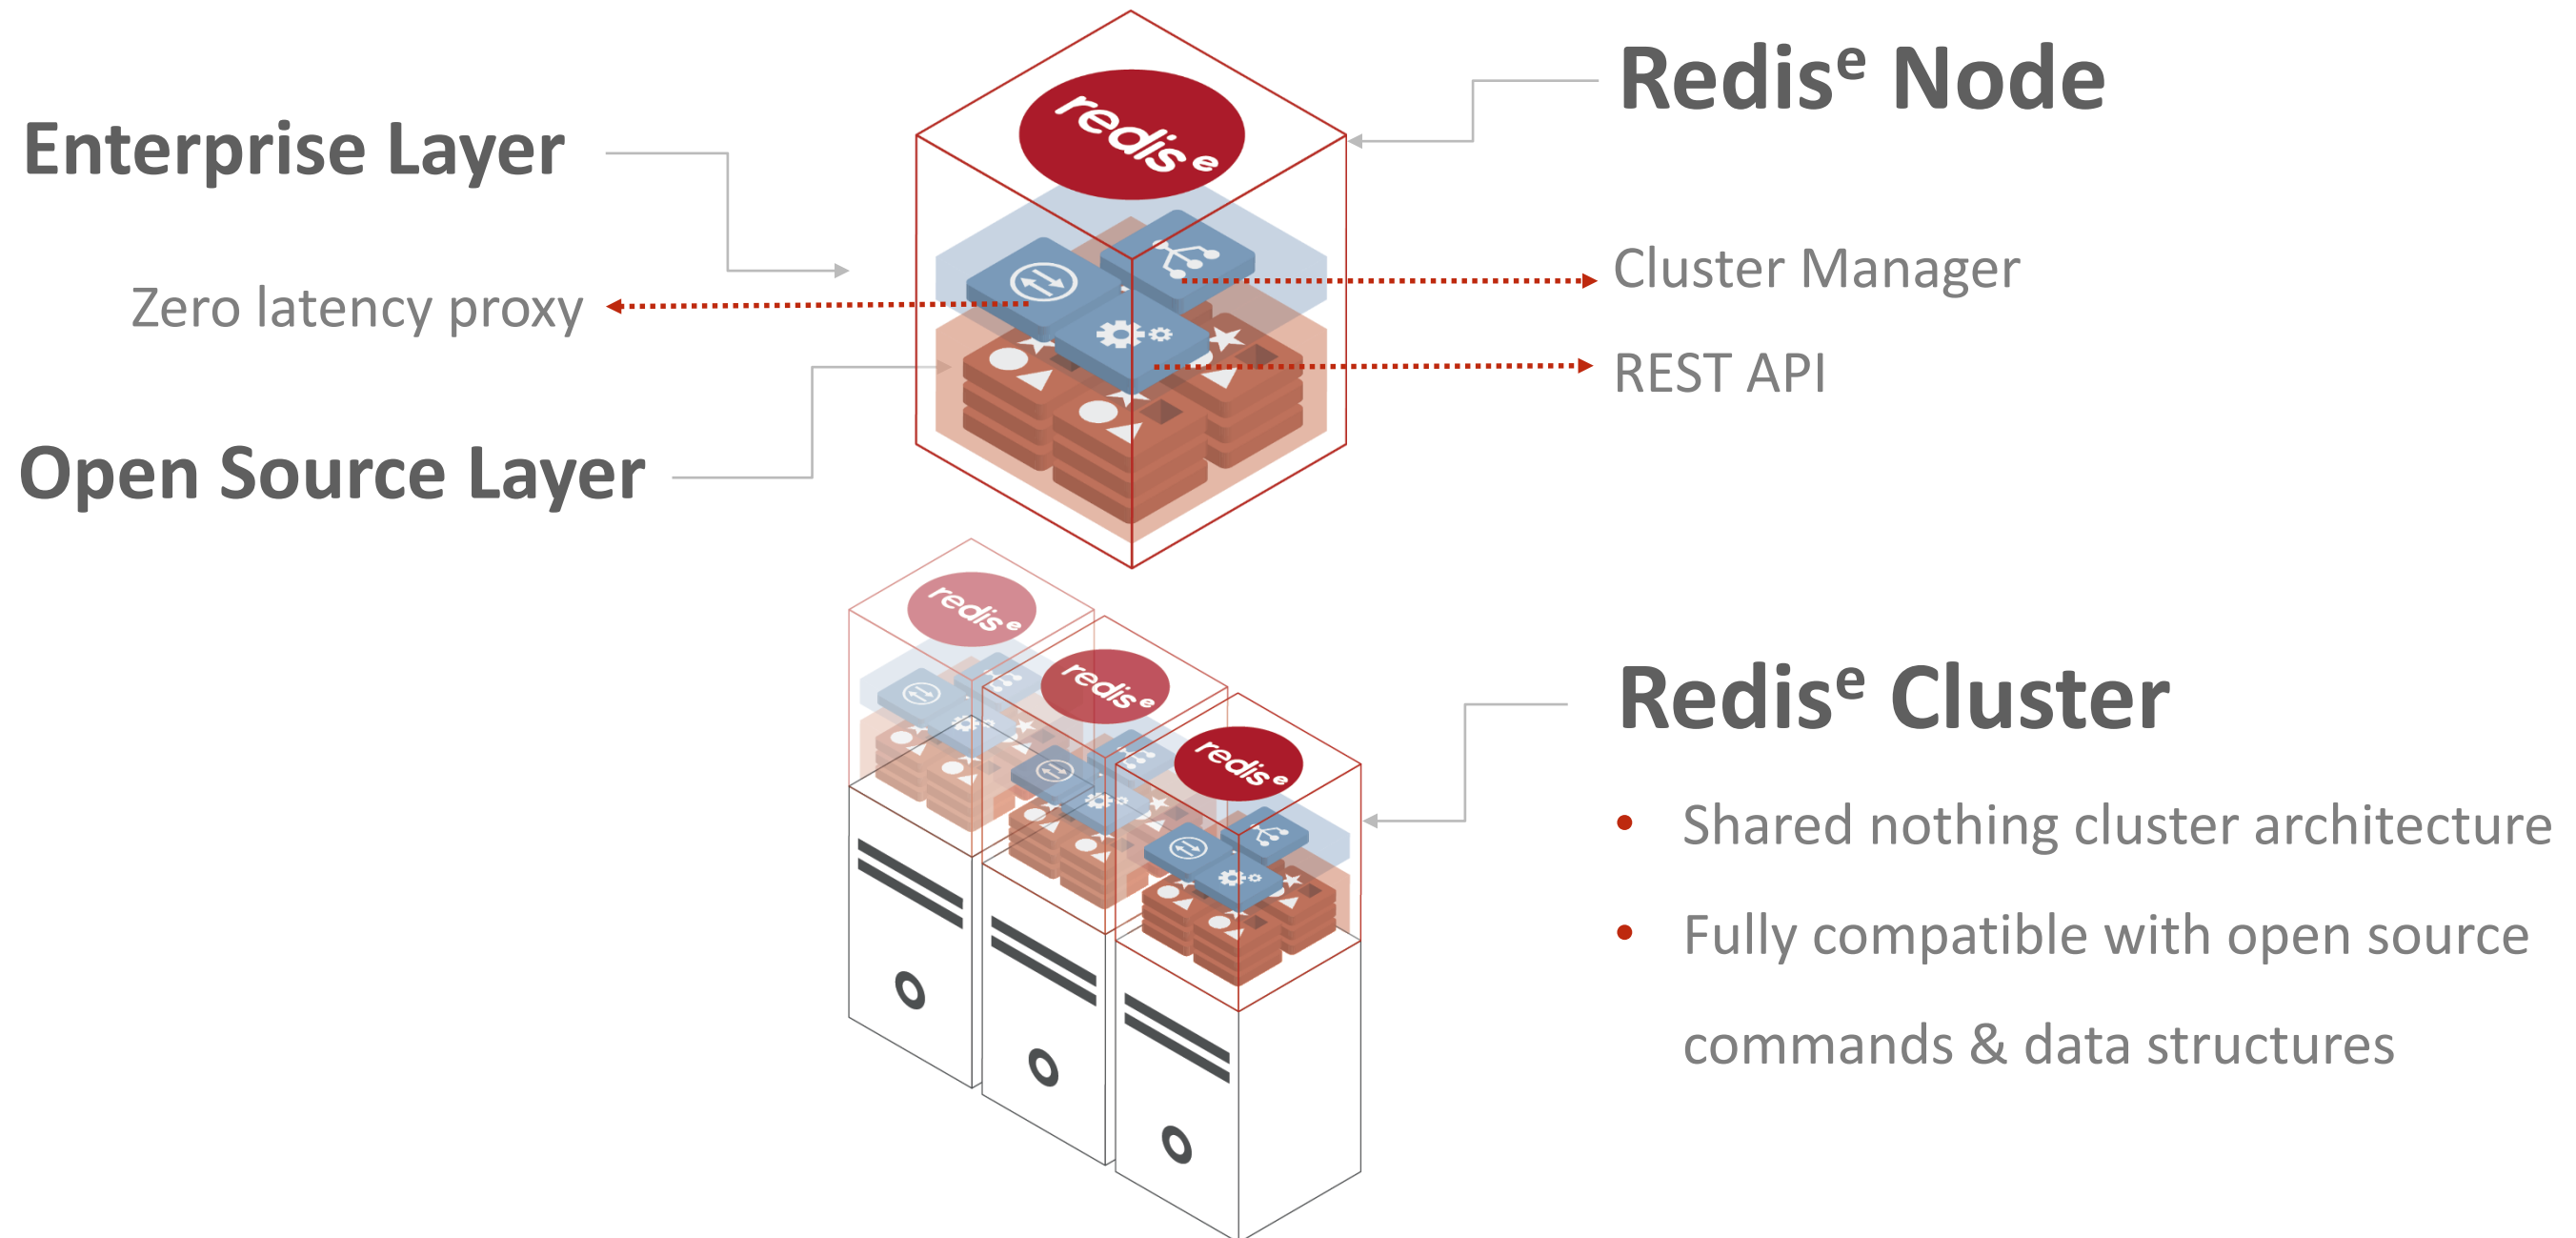
High Availability with Redis^e

- Replication Architecture in Redis^e
 - Local Replication across the LAN
 - Cross-Geo Replication across the WAN

Architecture Overview


Redis Architecture

Redis^e - Open Source & Proprietary Technology



Redis^e - Shared Nothing Symmetric Architecture


Distributed Proxies
Single or Multiple Endpoints



Unique multi-tenant “Docker” like architecture enables running hundreds of databases over a single, average cloud instance without performance degradation and with maximum security provisions

Redis^e - Multi-Tenancy

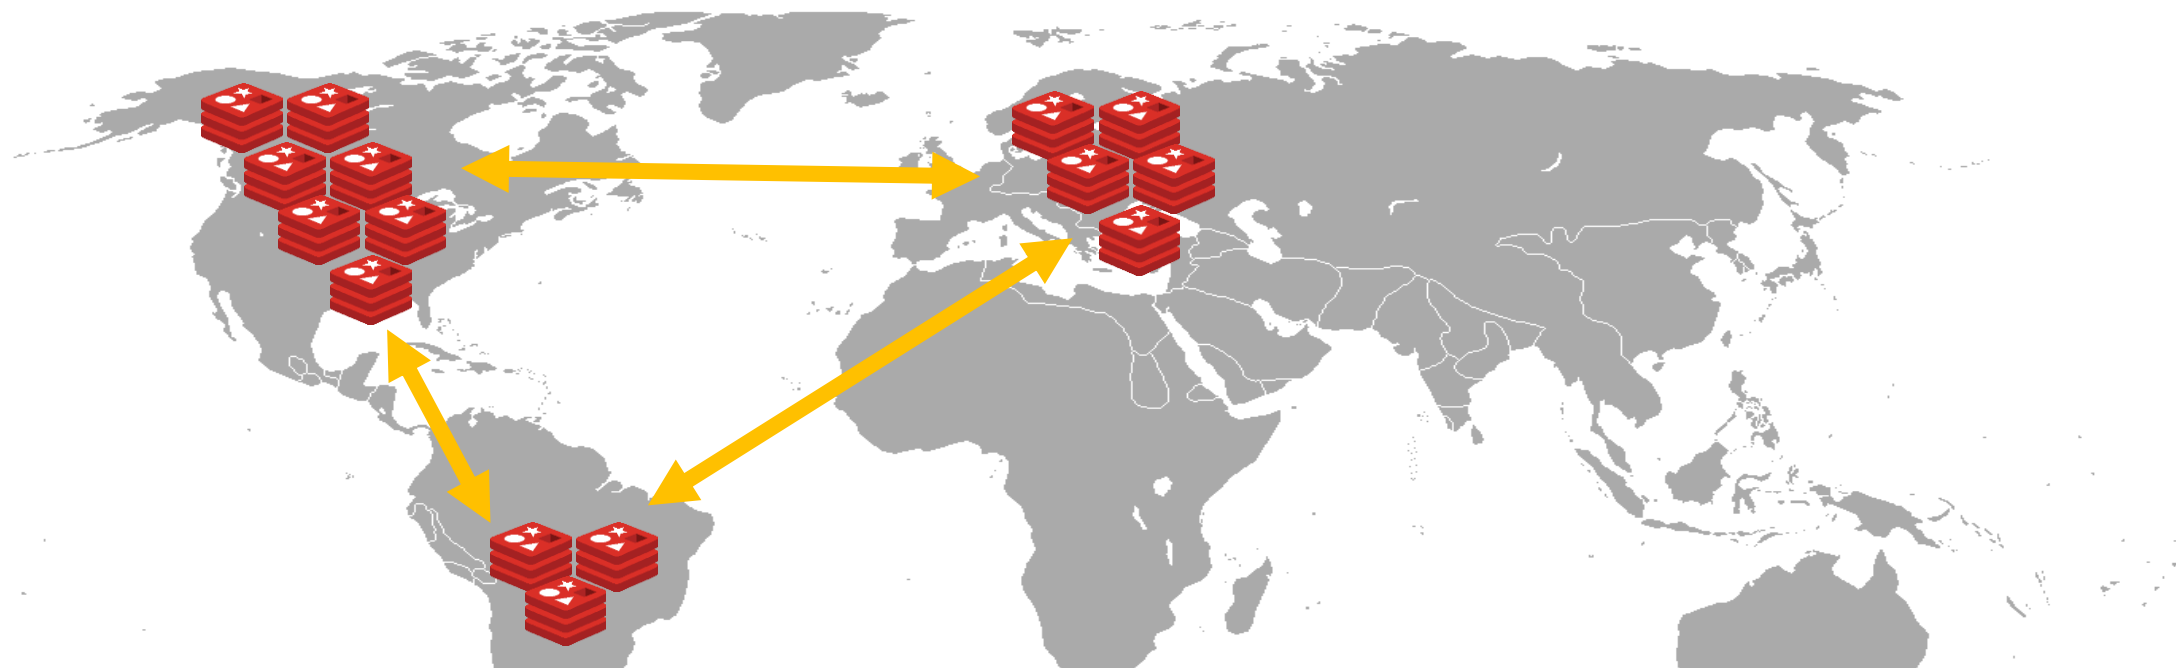
200+ customers/shards on a single 4vcore cloud instance



- Shard isolation/protection
- Noisy-neighbor cancellation
- Minimizing CPU consumption of inactive customers


Redis^e : State-of-the Art Database

- Multi-model, can support all popular database models and modern use cases
- ACID compliant with support for multi command/operation transactions
- Geographically distributed, “active-active”, multi-master architecture, with “strong eventual consistency”, based on CRDT (Conflict-free Replication Data Type) technology



Redis^e Benefits


Effortless Scaling



Simple, Seamless
Clustering. Linear scaling



ACID Compliance in
Cluster Architecture



Cluster Support for
Redis Modules

Always On Availability



Instant Failure Recovery,
No Data loss




Stable and Predictable
High Performance



Active-Active Geo
Distribution

Substantially Lower Costs



Run on Flash as a RAM
extension



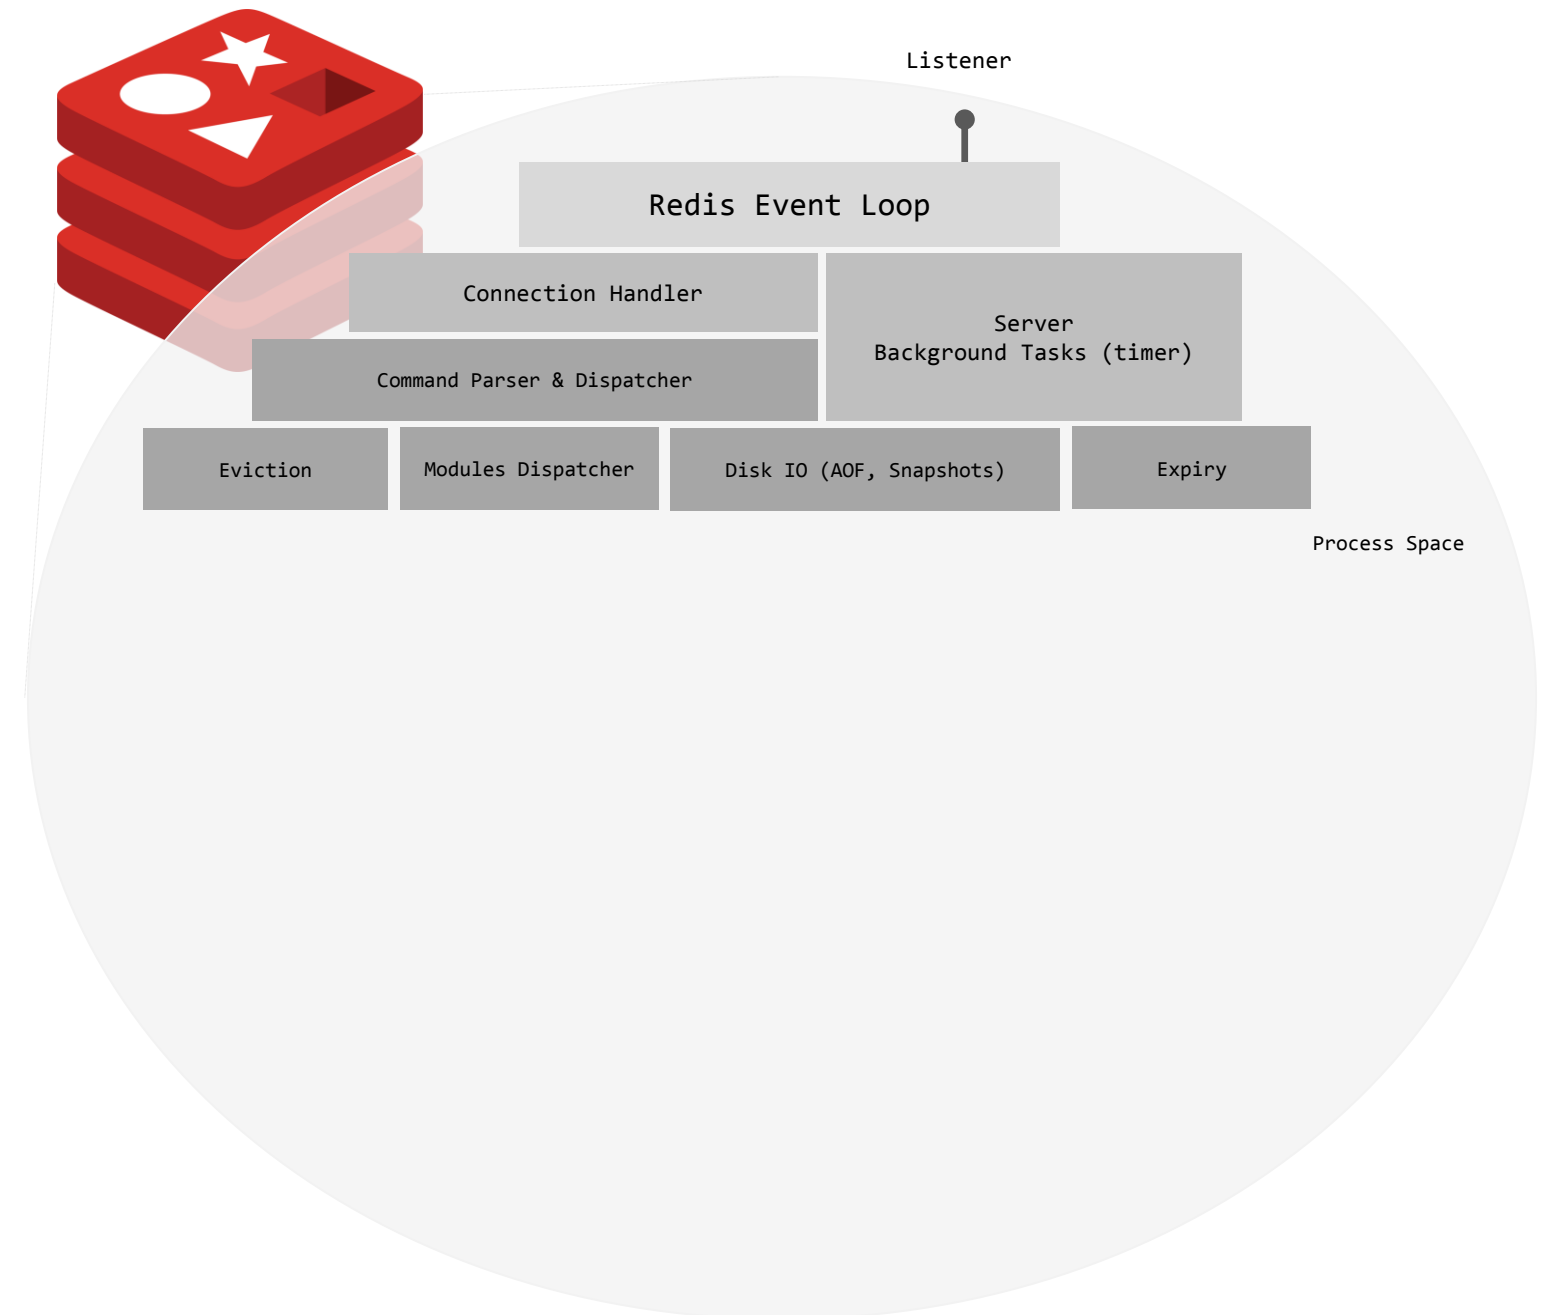
Fewer resources, lower
overhead




Top notch 24x7 expert
support

Redis Architecture

- Single Threaded, In-memory Engine with Persistence
- “Lock Free” architecture for fast execution



Redis^e Technology – Node Architecture



Redis^e Technology – Cluster Architecture



Redis^e Cluster

- Shared nothing cluster architecture
 - Single node type for simple scalability
- Fully compatible with open source commands & data structures
 - Simply change your connection string to Redis^e

DEMO

Quick Tour of the Redis^e Pack Cluster

Anatomy of Read & Write

Read/Write Operation with Redis^e

1. App submits the operation to one of the Proxies in Redis^e

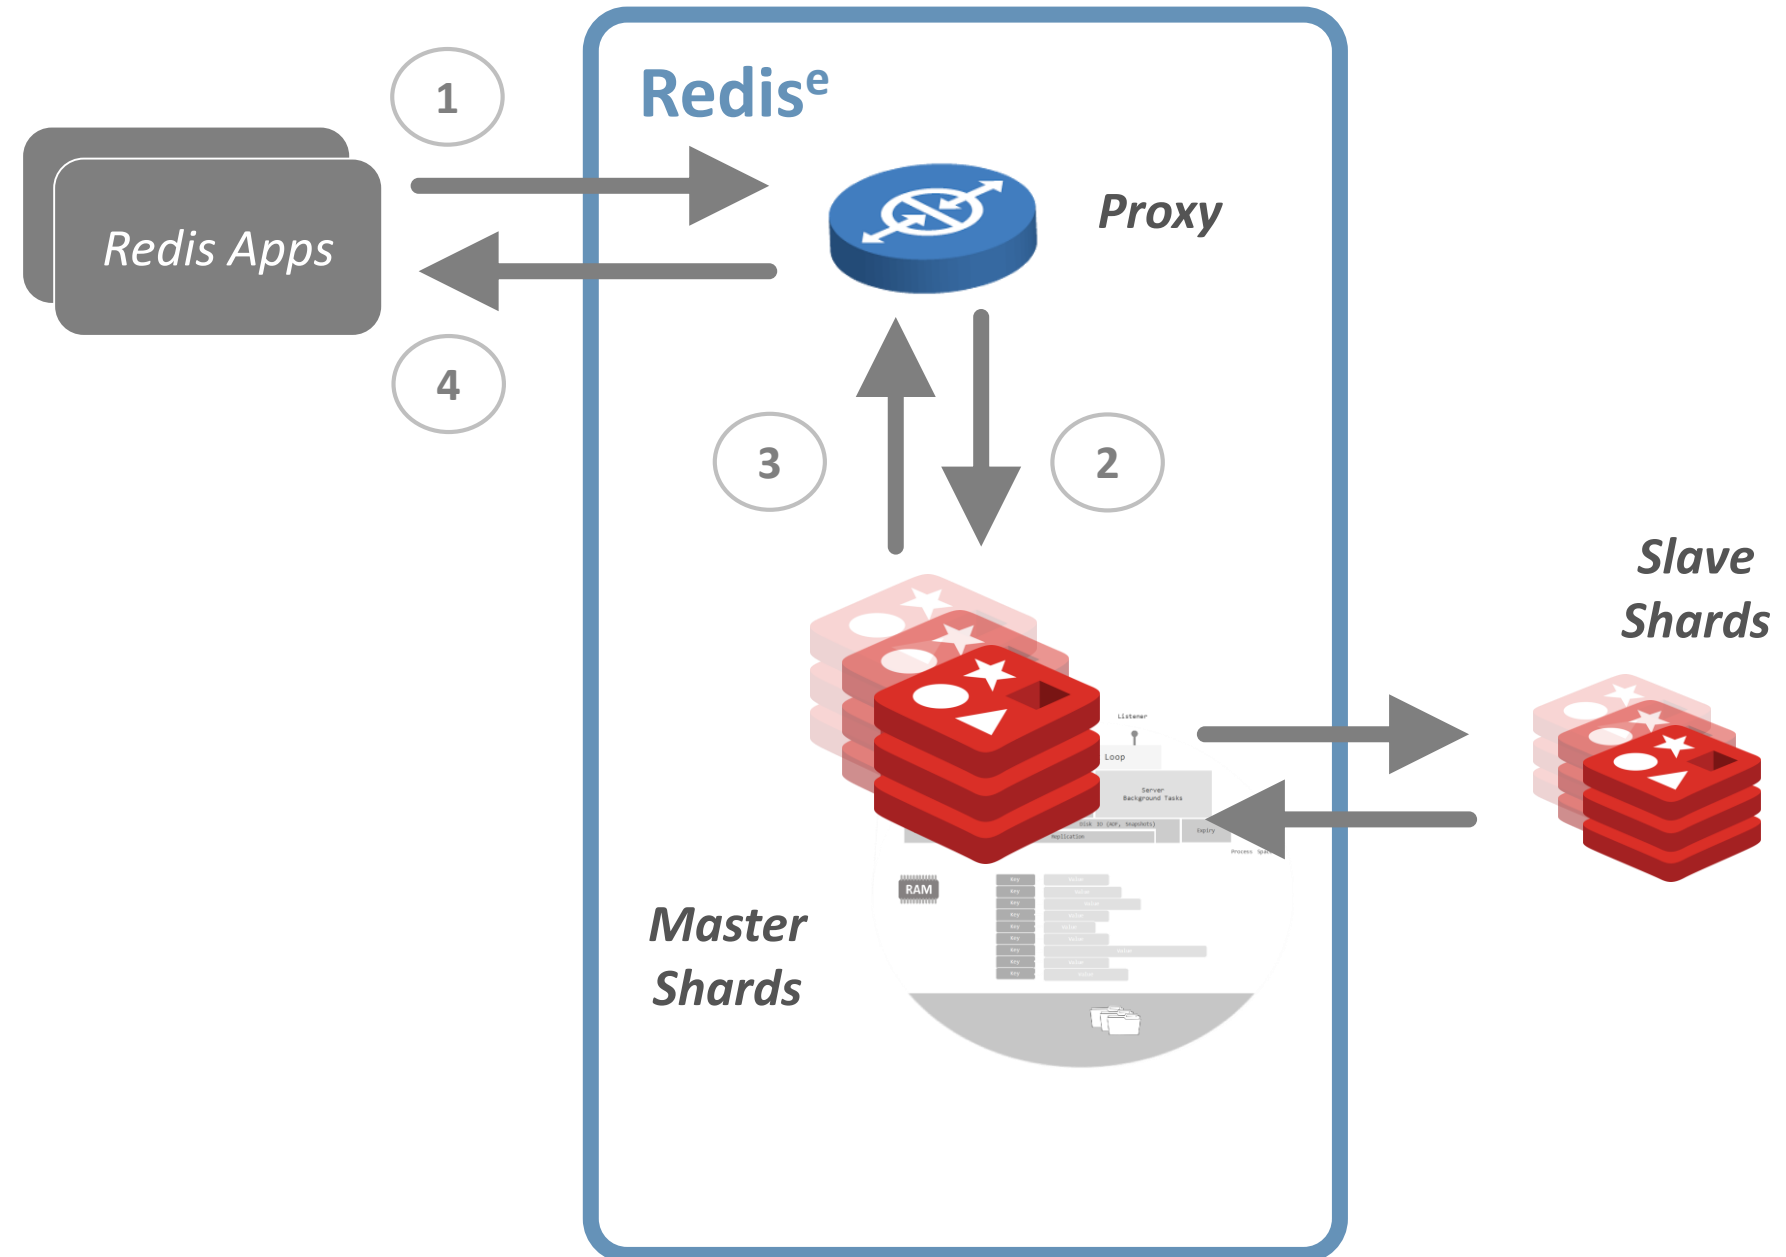
- Single Key Ops (*GET, STRLEN, HSTRLEN* etc)
- Multi Key Ops (*MGET, BRPOP, EXISTS, TOUCH*, etc)

2. Proxy distributes the operations to the corresponding shards in parallel

3. All involved shards return data to proxy

- a) Replicate writes to slave shards

4. Proxy assemble responses back to App



Read/Write Operation with Redis^e Flash

1. App submits the operation to one of the Proxies in Redis^e

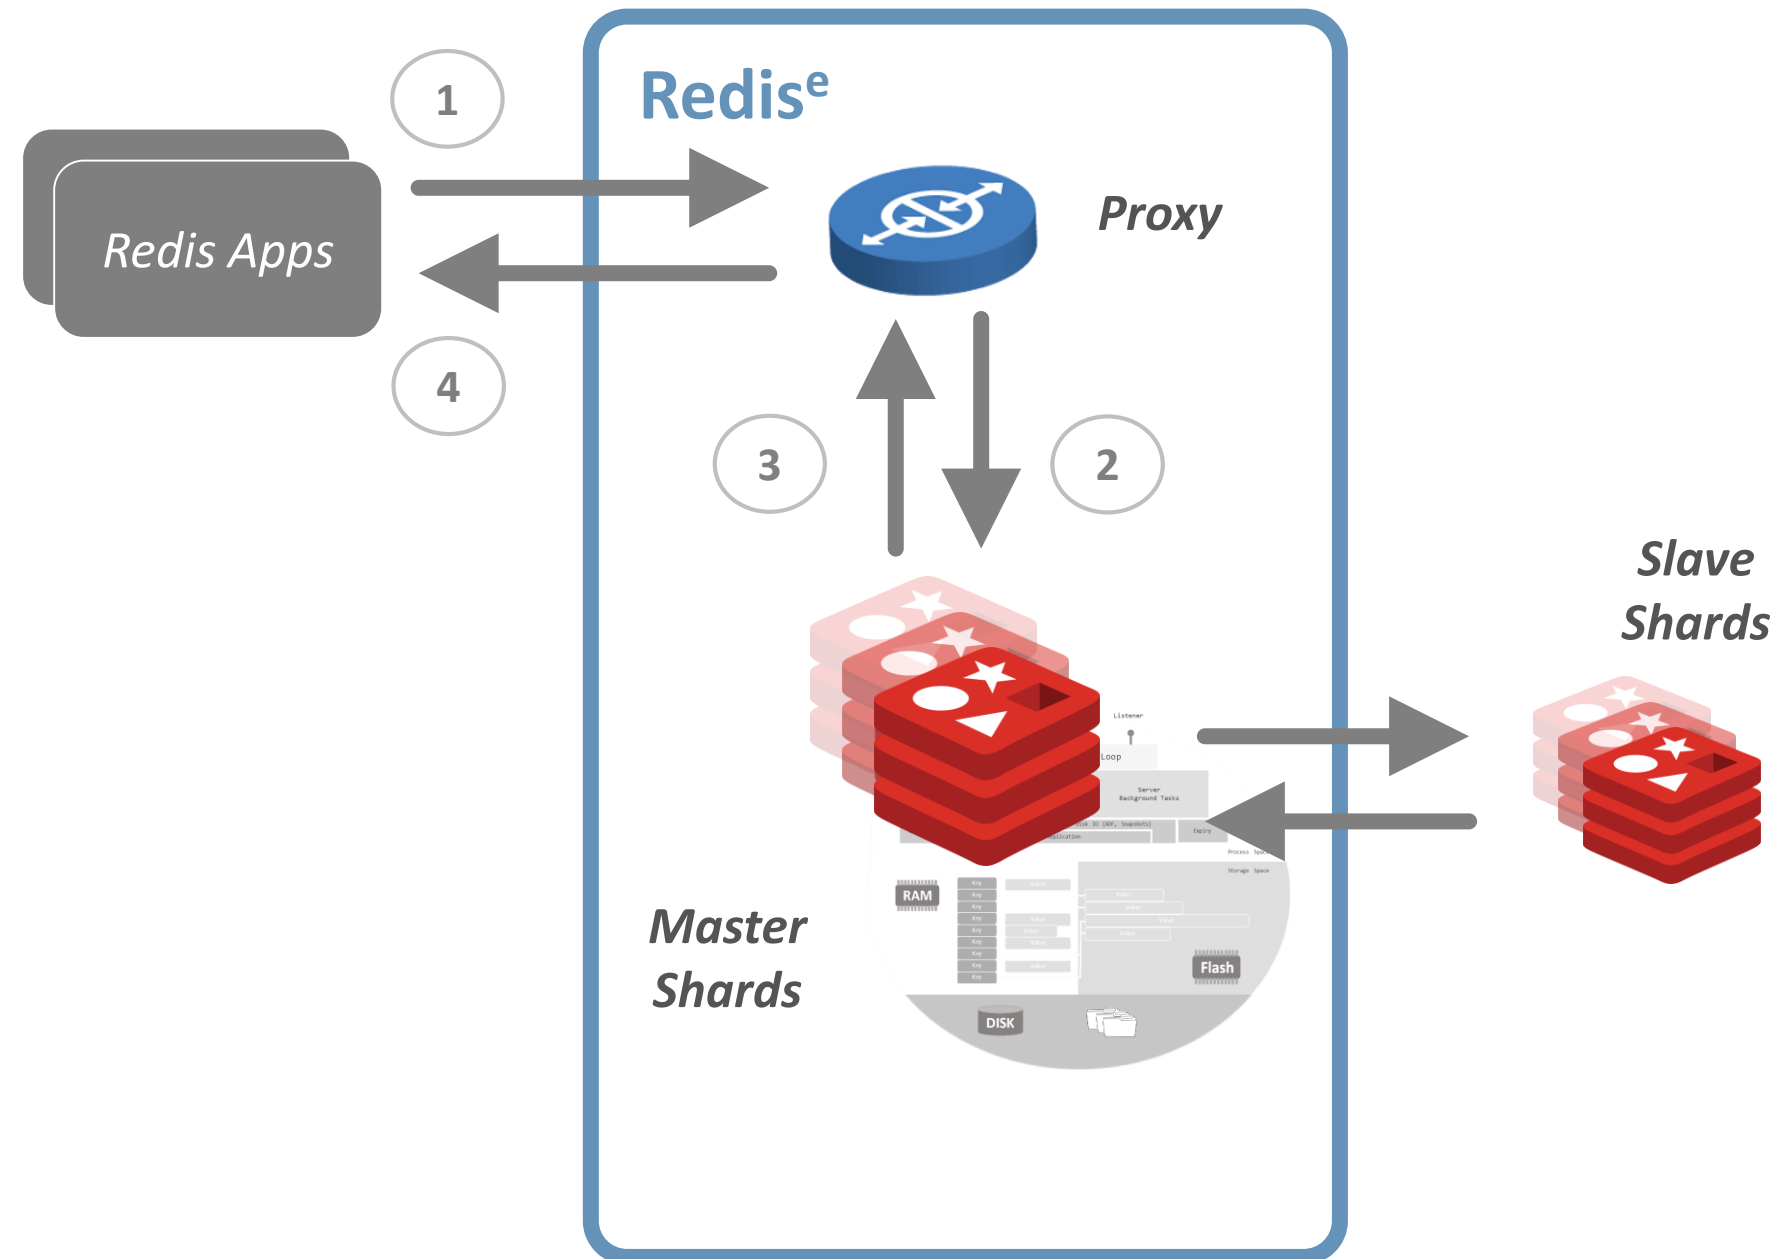
- Single Key Ops (GET, STRLEN, HSTRLEN etc)
- Multi Key Ops (MGET, BRPOP, EXISTS, TOUCH, etc)

2. Proxy distributes the operations to the corresponding shards in parallel

3. All involved shards return data to proxy

- a) Fetch value from flash if not in RAM
- b) Replicate writes to slave shards


4. Proxy assemble responses back to App



Scaling Applications with Redis^e

Scaling Throughput


Redis^e Shards



Database is sharded by the Cluster Manager.
Each shard runs on a separate process.

Scaling Up and Scaling Out


Rebalancing & Resharding without Downtime



Multi Tenancy

Built for Ground up for Multi-tenancy

- Simple Tenant Manageability,
- Tenancy Model for the Best Economics.
 - Mixing Many Small Tenants & Isolated Large Tenants



How is Scaling Managed

Redis^e Cloud

- Fully-automated

Redis^e Pack


- GUI
- CLI
- API

Scaling Data Size


Redis^e Flash

Massive Datasets with Near-Ram Latency at a Drastically Lower Cost

- Optimized Read/Writes with RAM-Extension approach
- Gain speed with smart caching between RAM and Flash
- Future proof for upcoming *persisted-memory* technology




Price/Performance – Memory Technology



Why Redis^e Flash?

Lower cost for Larger data sets

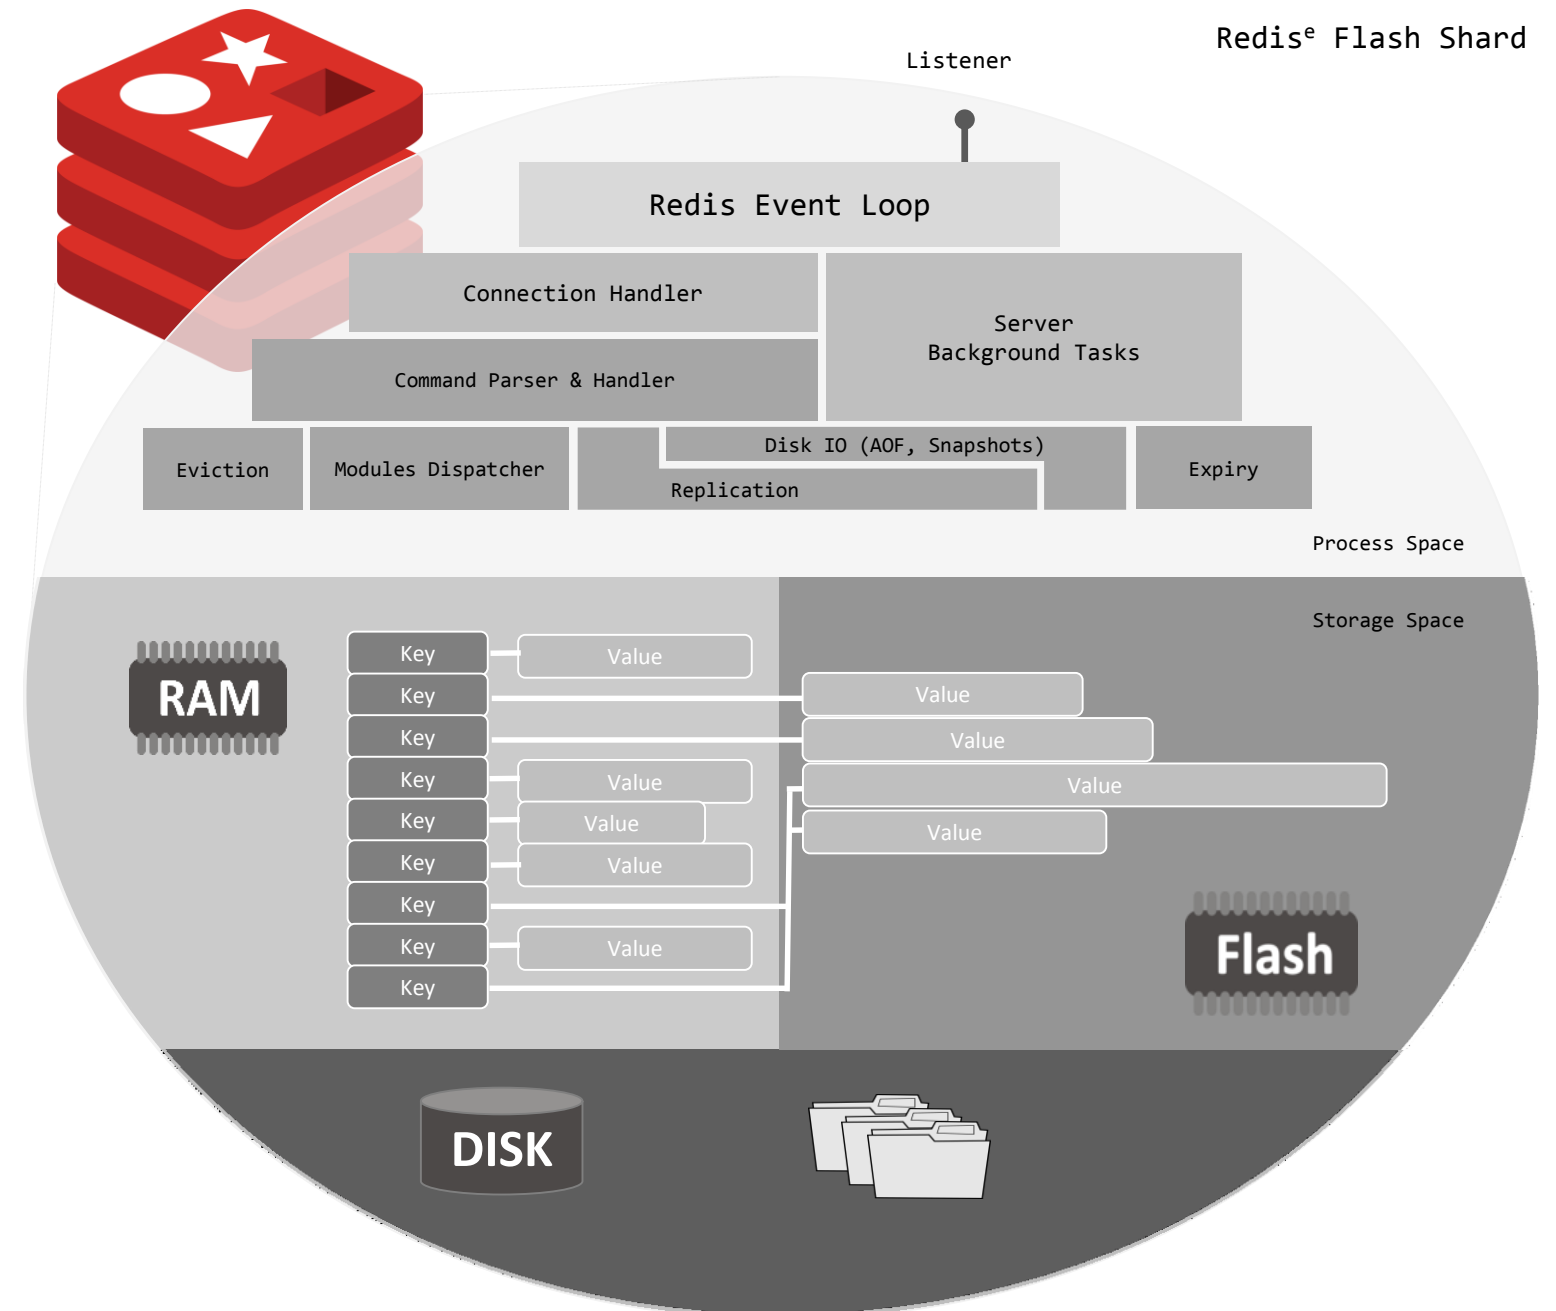


>80%

Lower Cost with RAM+Flash
Compared to all-in-RAM

Redis Shards in Redis^e

- Redis^e Flash Shard
- Ability to extend RAM to Flash for cheaper storage of data (Redis^e Flash)
- Advanced threading for better parallelism



Redis^e Flash vs Disk Based Databases?

	Redis on Flash	Disk Based Databases
Hot Value Handling	No IO Required Keep hot values in RAM	Heavy IO Required Keeps writing to disk
Write Performance	Faster Writes Non-Durable Writes with Ram Extension approach	Slower Writes Durable Writes (WAL, Redo logs etc)
Cloud Optimized	Fast Local Writes to Ephemeral Drive Utilizes the Ephemeral Drive for fast local IO and Network IO for durability	Slow Writes to Network Attached Storage <i>CANNOT</i> Utilizes the Ephemeral Drive for fast local IO
Future Proof	Ready for Persistent Memory Systems like Intel 3D-XPoint	Needs Re-Architecting



High Availability with Redis^e

Redis^e Replication Architecture

Replication with Redis^e



Low Replication Lag & High Replication Throughput

- Local Replication: Built for LAN
 - Higher bandwidth
 - Lower latency
 - High quality links susceptible to fewer failures and retransmits
- Cross-Geo Replication: Built for WAN
 - Lower bandwidth
 - Higher latency
 - “Noisier” network quality susceptible to more failures and retransmits



Cross-Geo Replication

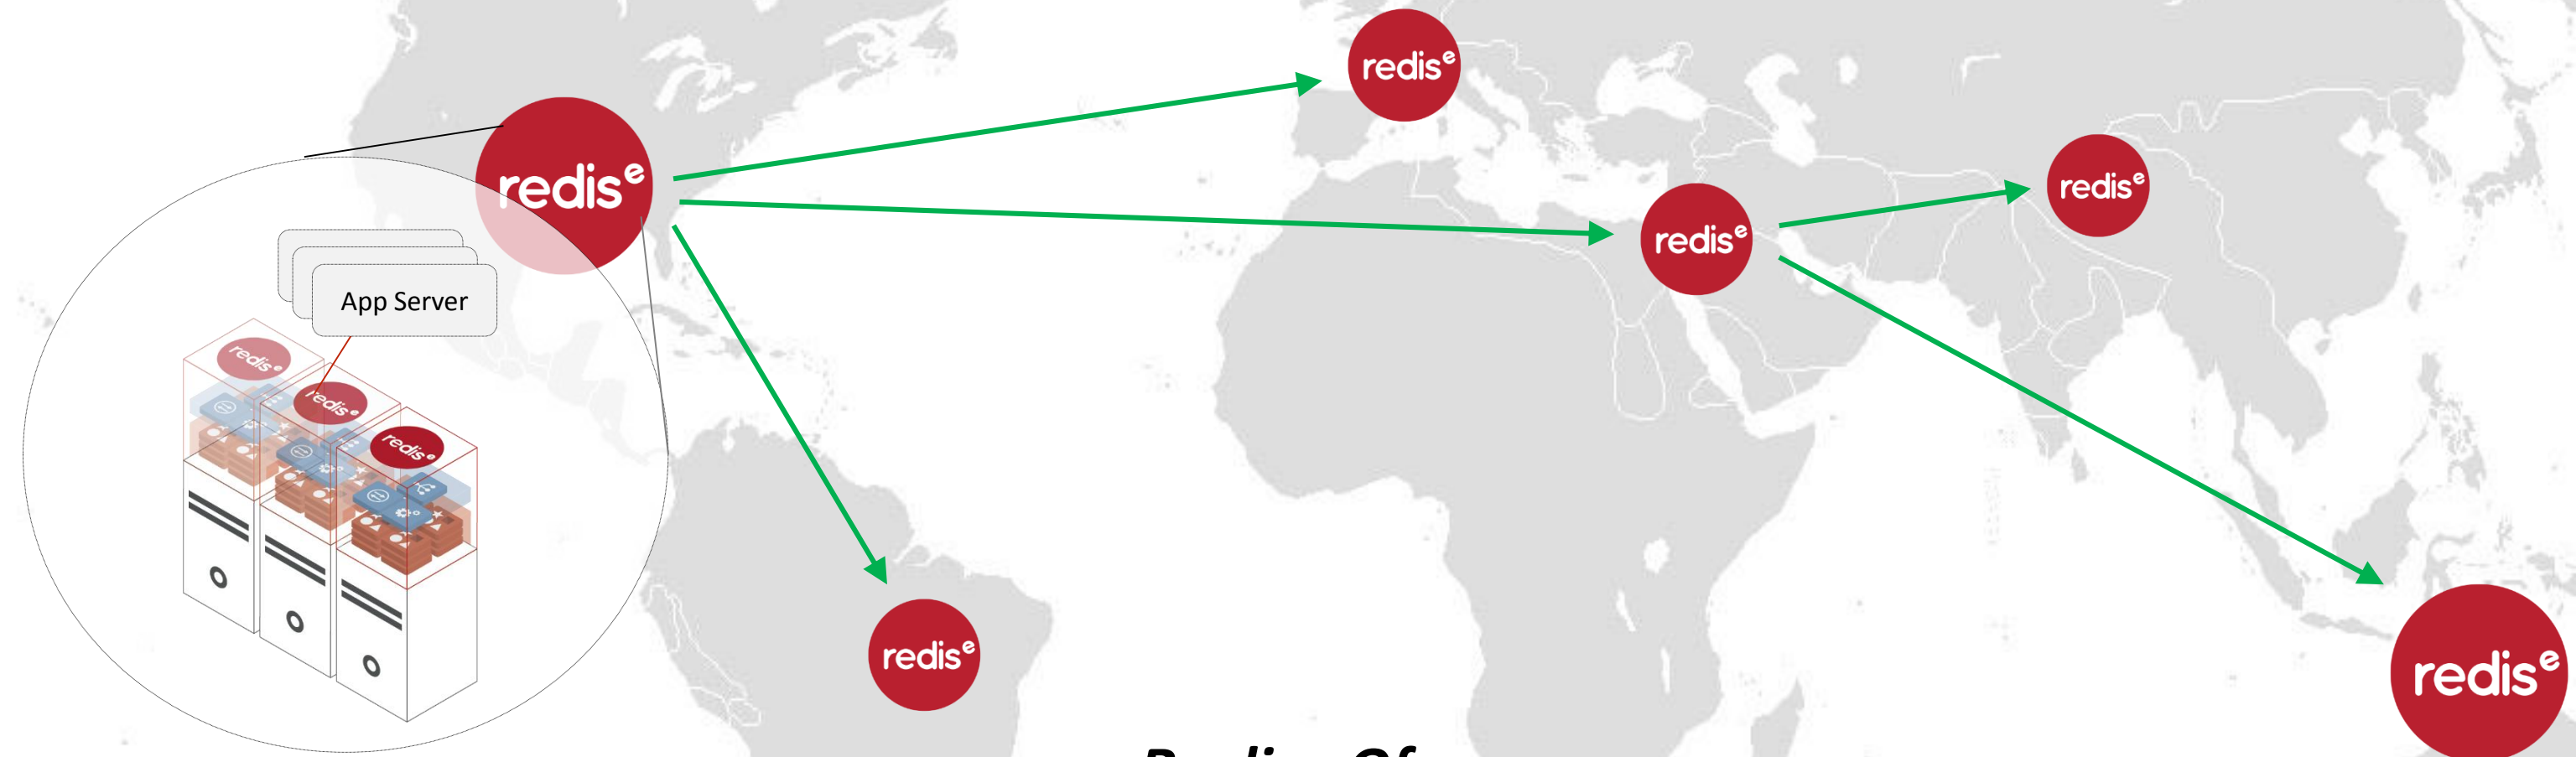
- Unidirectional Replication
 - **Replica Of** - Source DB to Destination DB across WAN
 - Content Distribution to Various Geographies for low latency, local reads
 - Geo-distribution for Hot Standby protecting against regional failures
- Multi-Master Replication *(coming soon in v5.0)*
 - **Multi-Master Databases** with Concurrent Active-Active Writes using CRDTs
 - Database spanning multiple geographies and clusters
 - Smart handling of conflicting concurrent writes with strong eventual consistency



Common Redis^e Geo Distribution Topologies

Topology #1

Geo Distribution for Fast Local Data Access



Replica Of
Geo Distribution for Local Data Access (CDN Like)

Great for...

Distribute Content Close to Your Users

- Read local copy with low latency, instead of crossing borders
- Push updates to all regions with fast, memory based replication

Avoid Write Conflicts - Achieve Higher Consistency

- All writes go to the master cluster, source database

Increase Read-Availability through Multiple Read-Replicas

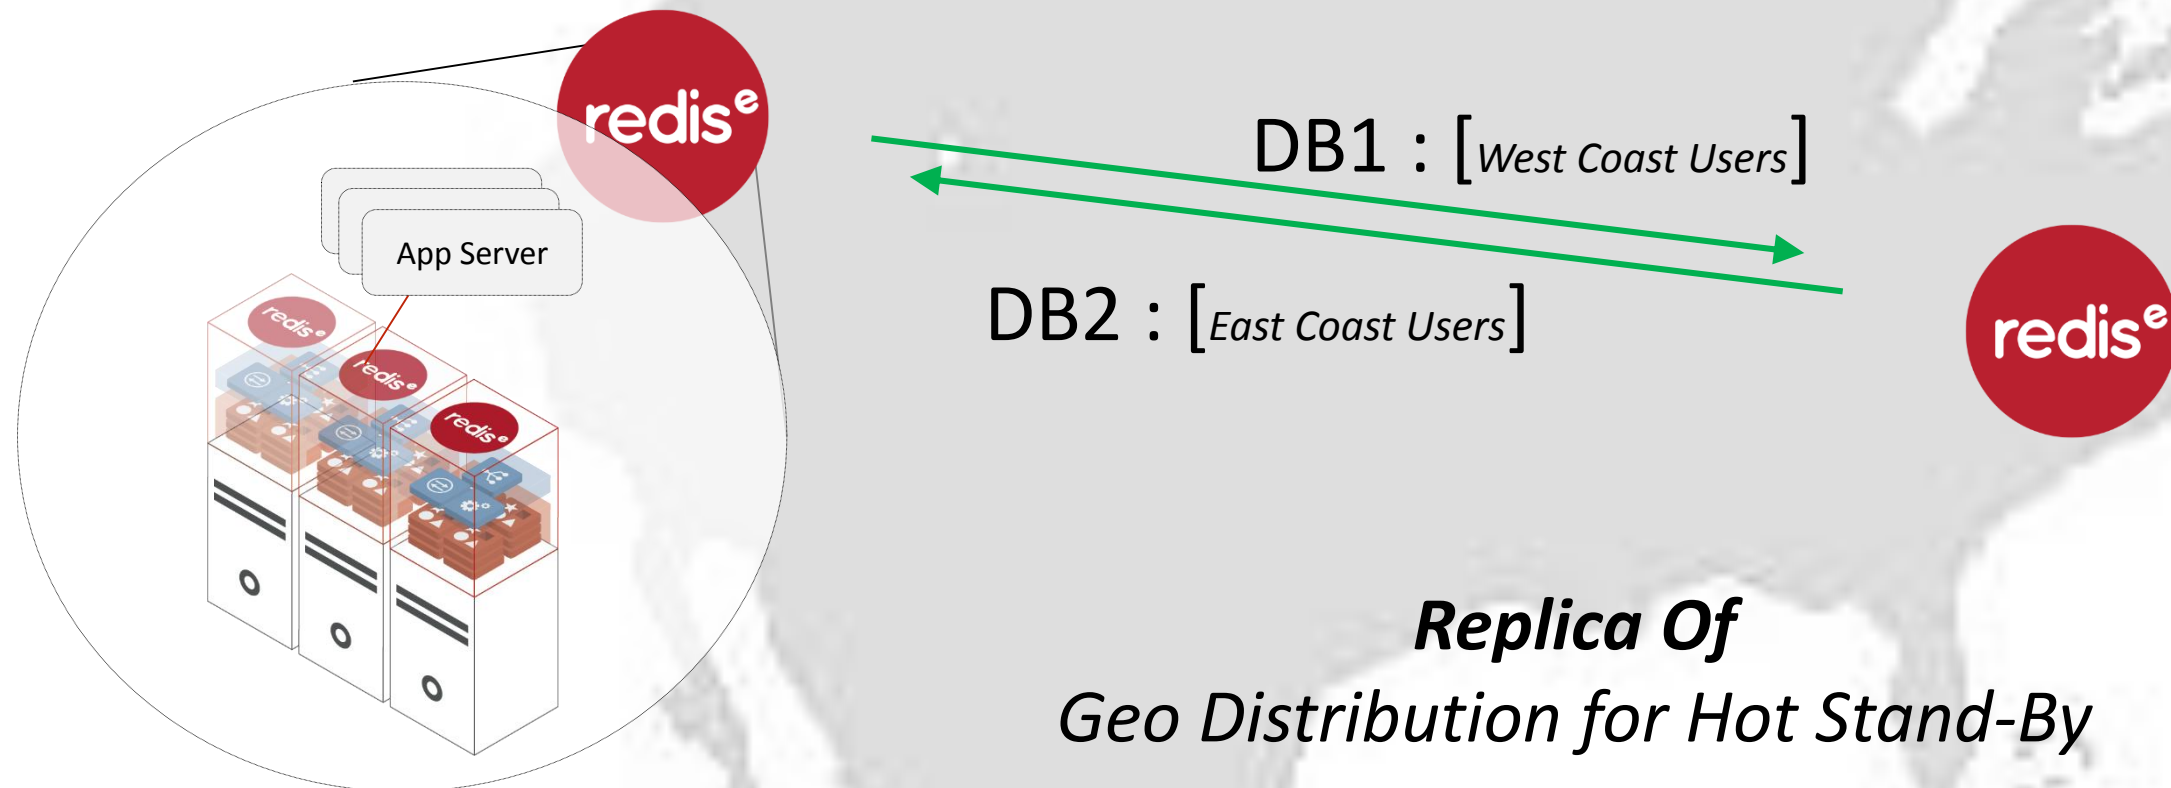
- Multiple data sources to read from maintained up to date with fast replication

Independently Size Each Region's Cluster

- Independently build enough capacity to handle local workloads in each region

Topology #2

Geo Distribution for Warm Stand-by



Great for...

Protect against Regional Failures

- Place a hot standby server to a different geo
- Increase Read and Write-Availability through Active-Passive Distribution

Avoid Write Conflicts - Achieve Higher Consistency

- All writes go to the master cluster with the source database

Independently Size Each Region's Cluster

- Independently build enough capacity to handle local workloads in each region

Recap

Redis and Redis^e (Redis Enterprise) Architecture

- Shards, Nodes, Clusters
- Replication in Redis^e
- Anatomy of Read/Write Operation

Scaling Applications with Redis^e

- Scaling Throughput – Redis^e Resharding and Rebalancing
- Scaling Connections - Redis^e Proxy
- Scaling Data Size – Redis^e Flash

High Availability with Redis^e

- Replication Architecture in Redis^e
 - Local Replication across the LAN
 - Cross-Geo Replication across the WAN

DEMO

Performance and HA with Redis^e Pack Cluster